

The Trident G-Loop: A Unified Cognitive Architecture for Adaptive, Near-Critical Intelligence

Abstract

- One loop, eight phases, operating inside a **Ψ -band** (near-critical regime) by regulating two interacting axes: a **Cognitive Resilience axis** (difficulty set-point F^* and precision-weighted gap $\Delta \hat{F}_t = \kappa \cdot (\mathcal{E}_t - F^*)$) and an **Inference axis** (precision-bias $b = \beta_d(d - d^*)$ and temperature $T(\chi)$).
- **Trident geometry**: subcritical **Autopilot shaft** \rightarrow near-critical **branch point** ($\mathcal{E} \approx F^*$) \rightarrow **Control** (compress/exploit) vs **Creative** (decompress/explore) prongs, coordinated by **Salience** (monitor/orient/reset).
- Two-sentence coda on the **Ψ -dynamic range of \mathcal{G}** ($\mathcal{G} = r \times G_f$): η -driven F^* ratcheting and controllable $b/T/\lambda$ spread; list telemetry: **Ω_Ψ** (band occupancy), **$\tau^{1/2}$** (recovery half-life), **H_Ψ** (hysteresis), **$\mu_{\{F^*\}}$** (set-point mobility).
- Reasoning in Trident G is realised by two cooperating WM workspaces: a relational, map-based LT-WM (FPCN-A \leftrightarrow hippocampus/DMN) and a rule/sequence ST-WM (FPCN-B \leftrightarrow DAN), with salience arbitrating the hand-off inside the Ψ -band.

1. Introduction

1.1 The challenge of flexible intelligence (stability \leftrightarrow flexibility).

1.2 The **Ψ -band principle**: benefits of operating between order and chaos; band (not point) via hierarchical/heterogeneous networks.

1.3 The **G-Loop & Trident metaphor** (Autopilot shaft, branch, two prongs, salience hilt).

1.4 The **Two Axes**: **Resilience** $\{F^*, \Delta \hat{F}_t = \kappa(\mathcal{E}_t - F^*)\}$ and **Inference** $\{b = \beta_d(d - d^*), T(\chi), \lambda\}$; preview **η** (competence/recovery) and **χ** (inconsistency/volatility).

1.5 Evolutionary rationale. We view human reasoning as an extension of Phase-3 mapping and Phases-4/5 execution: **abduction + counterfactuals** enable safe search in uncertain niches, while **deduction + induction** compress and transmit validated structure. Salience arbitrates these modes to keep operation near **Ψ** , yielding adaptive intelligence in variable environments.

1.6 Roadmap.

2. Theoretical Foundations & Core Commitments

2.1 **Policy objectives:** neutral score $J(\mathbf{a}|\mathbf{s})$; Active Inference uses $J = -G_{\{\mathrm{EFE}\}}$, classical decision uses $J = \mathrm{EV}(\mathbf{a}|\mathbf{s})$; selection via $\pi(\mathbf{a}|\mathbf{s}) \propto \exp(J/T(\chi))$; precision control via \mathbf{b} and $T(\chi)$.

- Use **CVaR_q** (e.g., $q=0.95$) or a **τ -quantile** value function when tail-shape telemetry (Sec. 6) flags heavy-tail/volatility regimes.
- Tie the switch explicitly to telemetry thresholds (ζ /J-ratio/ES from Sec. 6) so readers see the closed loop: “The risk-sensitive objective engages when $\zeta < \zeta^*$ or J-ratio $> J^*$ or $ES_q > ES^*$ ”

2.2 **Relational core (ϕ):** hippocampal–prefrontal **predictive map** (SR-style), value-landscaped by **EV** or $-G_{\{\mathrm{EFE}\}}$; learned translators to modalities.

2.3 **Metacognition as cross-loop consistency:** $\chi_{\{\mathrm{meta}\}}$ monitors inter-loop inconsistency; brief monitor pulses to re-centre.

2.4 **Multi-timescale architecture:** fast/medium/slow controllers; rhythms ($\psi/\theta/\delta$) as dynamic weights (non-exclusive).

3. Controllers & State Variables (the meta-cognitive toolkit)

3.1 **State signals:** η (competence/recovery), χ (uncertainty/inconsistency).

3.2 **Controllers:**

- **F*** thermostat (Resilience axis; branch set-point).
- \mathbf{d}, \mathbf{d}^* (dopamine tilt & neutral set-point) $\rightarrow \mathbf{b} = \beta_{\mathbf{d}}(\mathbf{d} - \mathbf{d}^*)$ (Inference tilt).
- **T(χ)** (policy temperature), λ, λ^* (generalisation width & target), α (plasticity), κ (gate gain in $\Delta\hat{\mathbf{F}}$), σ (reasoning budget), \mathbf{v} (decompression/novelty budget).

3.3 $\eta \rightarrow \mathbf{F}^*$ coupling: competence-driven upward drift (auto-difficulty).

4. Core Hypotheses (with short tests)

H1 — SOC-for-Inference. Homeostatic plasticity drives hierarchical networks into a **Ψ -band** that maximises information capacity/dynamic range required for good policy performance; loops track **F*** and then choose policies by J (EV or $-G_{\{\mathrm{EFE}\}}$). *Test:* shift E:I toward/away from balance and track Ψ -metrics + performance/ J -proxies.

H2 — Trident Branching. Near $\mathcal{E} \approx \mathbf{F}^*$ and elevated $\chi, \Delta\hat{\mathbf{F}}$ and salience arbitrate a bifurcation

into **Control** (compress) vs **Creative** (decompress), with monitor pulses to re-centre **b/T**. *Test*: SN-led switches into FPCN-B (Control) vs FPCN-A/DMN (Creative) with prong-specific telemetry.

H3 — Global + Per-loop Ψ -bands. Multiple G-loops operate with loop-specific Ψ_k windows whose intersection defines a global Ψ ; $\chi_{\{\text{meta}\}}$ triggers monitoring when loops disagree. *Test*: show Griffiths-like banding across scales and salience-preceded re-entry.

H4 — Ψ -Dynamic-Range of \mathcal{G} . $\mathcal{G} = r \times G_f$ reflects (i) **Resilience range** (tolerable $|\mathcal{E}-F^*|$ and η -driven $\mu_{\{F^*\}}$) and (ii) **Inference range** (controllable **b/T/** λ span with clean re-entry), aggregated across loops. *Test*: derive a Ψ -profile (Ω_Ψ , $\tau_{1/2}$, H_Ψ , $\mu_{\{F^*\}}$) and relate to G_f/G_c ; show training gains.

H5 — Two Creative Routes. Exploration enters via (i) boredom/under-challenge (soft D2) and (ii) error/volatility (hard D2), with distinct arousal/ACC signatures. *Test*: dissociate behavioural/physiological markers for the two routes.

H6 — Dual WM Workspaces. FPCN-A \leftrightarrow hippocampus supports long-horizon relational “map” (Creative/Plan), FPCN-B \leftrightarrow DAN supports rule/sequence execution (Control), with salience arbitration. *Test*: mode-specific coupling shifts across phases.

5. The Eight Phases of the G-Loop (navigating the Trident)

1. **Sense & Move** (estimate \mathcal{E}_t).
2. **Gap-check & Mode Branch** (compute $\Delta \hat{F}_t = \kappa \cdot (\mathcal{E}_t - F^*)$; arbitrate via salience using χ , $\chi_{\{\text{meta}\}}$).
3. **(De)Compress ϕ** (Control: compress $\rightarrow \eta \uparrow$; Creative: decompress under $\lambda \downarrow / \nu \uparrow$).
4. **Decide/Gate** (softmax over \mathcal{J} with $\mathcal{J} = \mathbf{EV}$ or $-\mathbf{G}_{\{\text{EFE}\}}$; $\mathbf{T}(\chi)$ sets exploration; F^* does not enter softmax).
5. **Execute** (policies expected to reduce χ and align $\mathcal{E} \rightarrow F^*$).
6. **Reflect/Update** (adjust F^* , **b**, **T**; update η , χ).
7. **Consolidate (λ)** (propagate validated structure; quarantine speculative bits).
8. **Reset** (restore Ψ -band rhythms; prep next cycle).

6. Telemetry, Predictions & Falsification

- **Band signatures:** avalanche scaling/dynamic range; metastability peaks near branch; rare-region effects (Griffiths-like).
- **Prong states:** SN-led switches → FPCN-B (Control) vs FPCN-A/DMN (Creative).
- **Two exploration routes:** boredom vs error/volatility with distinct arousal/ACC patterns.
- **Map (de)compression evidence:** SR-like updates; pattern separation/completion.
- **Ψ -profile metrics:** r , G_f , Ω_Ψ , $\tau_{1/2}$, H_Ψ , $\mu_{\{F^*\}}$. *Falsify where moving toward near-criticality fails to improve Ψ -metrics or where SN activity does not precede re-entry.*

Tail-shape telemetry.

ζ (tail index): Hill α_{Hill} for power-law-like tails (optionally report excess kurtosis κ_{ex}). **Lower $\alpha_{\text{Hill}} \Rightarrow$ heavier tail.**

J-ratio (one-big-jump index): $J = \max_t |\Delta F_t| / \sum_t |\Delta F_t|$ over the window; **higher J \Rightarrow single-shock dominance.**

ES_q(ΔF): Expected Shortfall of ΔF at level q (e.g., $q=0.95$).

Control law: If $\alpha_{\text{Hill}} < \zeta^*$, or $J > J^*$, or $\text{ES}_q > \text{ES}^*$, then **raise $T(\chi)$** (broader sampling), **lower λ** (sandbox Creative), trigger a **Monitoring pulse**, and recentre **b**; relax when metrics fall below thresholds.

Prediction: Tail alarms should precede **$T(\chi) \uparrow$** , **$\lambda \downarrow$** , **Monitoring \uparrow** and reduce **ES_q(ΔF)** and band exits.

7. Minimal Computational Program (plus ablations)

Hybrid instantiation.

Wilson–Cowan E/I nodes (local Hopf-edge control) with Kuramoto-style phase coupling on a hierarchical-modular / empirical connectome.

Controller mapping.

$F^*F^*F^*$: slow homeostasis to target partial synchrony; bbb : subnetwork gain/tilt (via $d-d*d-d^{\wedge}*d-d^*$); $T(\chi)T(\chi)T(\chi)$: input-noise / frequency-jitter gate; $\lambda/\alpha/\lambda/\alpha$: propagation & plasticity knobs; J : EV/-G_{EFE} or **risk-sensitive (CVaR/quantile) when tail flags fire.**

Readouts (additions in bold).

- Global/meso order parameter $R(t)$; power-law exponents for (de)synchronisation; E:I telemetry; recovery $E \rightarrow F^* \rightarrow E$ and $\tau_{1/2}$ after shocks.
- **Tail-shape metrics on performance residuals ΔF :**
 $\zeta = (\alpha_{\text{Hill}}, \kappa_{\text{ex}})$; **J-ratio**
 $J = \max_t |\Delta F_t| / \sum_t |\Delta F_t|$; $J = \max_t |\Delta F_t| / \sum_t |\Delta F_t|$; **ES_q(ΔF)** (e.g., $q=0.95$).
- Gate & safety telemetry: $T(\chi)$, λ , **Ψ -band exits** (count/length).

Protocols & stress-tests.

1. **Baseline Gaussian (light-tail):** finite-mgf noise; expect CLT regime—risk-neutral J acceptable near mean.
2. **Heavy-tail / “one-big-jump”:** Pareto-like or α -stable noise + rare shocks; expect **risk-sensitive J** to reduce ES_q, band exits, and $\tau_{1/2}$.
3. **Injection tests:** single large outlier vs many small; verify **J-ratio** $\uparrow \Rightarrow T(\chi) \uparrow, \lambda \downarrow, T(\chi) \uparrow$, $\lambda \downarrow$ and faster safe re-entry.
Prediction: Only the **tail-aware** agent maintains high **Ψ -occupancy**, lower **ES_q(ΔF)**, fewer/shorter Ψ exits, and shorter $\tau_{1/2}$ under heavy tails.

Ablations.

- (i) **No $F^*F^*F^*$ thermostat** \rightarrow over/under-challenge; unstable $R(t)$; slow $E \rightarrow F^* \rightarrow E$.
- (ii) **No $\chi \rightarrow T(\chi) \rightarrow T$ gate** \rightarrow brittle exploration/exploitation; variance spikes; excess Ψ exits.
- (iii) **No η gate (monitoring pulse)** \rightarrow poor regime re-entry; longer $\tau_{1/2}$.
- (iv) **No risk-sensitive J** (force $EV/-G_{\{EFE\}}$) \rightarrow **catastrophic failure modes in heavy-tailed environments:** **ES_q(ΔF)** \uparrow , J-ratio events propagate, Ψ -occupancy falls, despite intact $F^*F^*F^*$ and $\chi \rightarrow T(\chi) \rightarrow T$.

Falsification criteria.

If enabling tail-aware J and tail telemetry **does not** (vs ablation iv) lower ES_q, shorten $\tau_{1/2}$, and reduce Ψ exits in the heavy-tail condition—claim is weakened.

8. Reasoning, Working Memory & Constraints

- **Deduction/Induction** = Control; **Abduction** = Creative; **Counterfactuals** at Decide/Act; **Analogy** bridges Creative→Control.
- **Constraint satisfaction**: feasibility gating/propagation (hard/soft/chance constraints) to guard Creative proposals and tighten Control propagation.
- **Dual WM**: long-term WM (FPCN-A↔hippocampus/DMN) vs short-term WM (FPCN-B↔DAN); salience hand-off. (Note - Both LT-WM (map) and ST-WM (rule/sequence) can operate in **either** creative (D2-biased) or control (D1-biased) mode; **D1:D2 sets processing style, not locus**—salience arbitrates the hand-off.)
- On this view, **reasoning is the control law** for staying inside the Ψ -band: Creative (abduction/analogy + counterfactuals) expands the hypothesis/policy space when χ is high; Control (deduction/induction + constraints) compresses and propagates validated structure for skillful action.”

9. Neuromodulatory & Rhythmic Implementation

- **DA** (D1/D2) → **d**, **d*** → **b** (stability–flexibility tilt); **LC–NE** → **T(χ)** (exploit–explore gain); **ACh** → precision/learning; **$\alpha/\theta/\gamma$** rhythms as dynamic weights.

10. Relation to Prior Work

- Active Inference, predictive coding, meta-RL, MDN/FPN accounts—**what’s shared vs what’s new**: the **F*** thermostat, **η/χ** gating, the **Ψ -band + Trident branching** geometry, and concrete, falsifiable telemetry.

11. Applications & Deployment Contract

- Learning/training, digital therapeutics, decision support; include a deployment table aligned to controllers ($\chi \rightarrow T$, F^* , η -gate, b , λ) and standard governance artefacts (model cards, risk registers, etc.).

12. Limitations & Ethics

- What's solid vs **partly supported/speculative** (e.g., equating dynamical criticality with an EFE set-point; global neutral D1:D2; χ/η as cross-loop constructs). State tests and caveats explicitly; clinical mappings as computational phenotypes, not diagnoses.
- **Supported:** near-critical banding / Griffiths-like effects; E:I tuning; salience-driven switching; $G_f \leftrightarrow MD/FPN$, $G_c \leftrightarrow DMN$ tendencies; metastability–creativity links; D1/D2 stability–flexibility trade-off; LC/NE as temperature-like.
- **Promising/speculative:** mapping dynamical criticality to an EFE-compatible set-regime; a global neutral D1:D2 (d^*); χ, η as cross-loop constructs; dual exploration routes within one loop.

13. Conclusion

- Re-state the **Trident** synthesis: Autopilot shaft \rightarrow near-critical branch \rightarrow prongs + salience hilt; two axes negotiating inside a **Ψ -band** to yield adaptive intelligence; concrete predictions and a compact simulation path.
-

Appendices

- **A. Mathematical formulae** (F^* thermostat, b updates via $d-d^*$, choice policy with J , λ controller; normal-form cartoon for branching).
 - **B. Glossary of variables** (η , χ , F^* , d/d^* , b , T , λ , α , κ , σ , v) and network labels (FPCN-A/B, DMN, DAN, SN).
-