



Adaptive intelligence balances stability and flexibility by operating in a narrow near-critical regime — the Ψ -band. We present the Trident G-Loop, a unified cognitive architecture that formalises processing in this regime as an eight-phase rhythmic cycle: Sense/Move → Branch → Map → Decide → Act → Update → Consolidate → Reset. Two interacting axes regulate the loop: a Cognitive Resilience axis (managing performance challenge via a difficulty thermostat, F^* , and the precision-weighted gap $\Delta \hat{F} = b \cdot (F - F^*)$) and an Inference axis (managing stability–flexibility via a D1:D2 precision-bias controller, b , together with the policy temperature $T(\chi)$). At the branch near $F \approx F^*$, a Monitoring/Salience coordinator arbitrates between Control (exploit/compress) and Creative (explore/decompress) modes, executing brief orient–reset–reconfigure actions when cross-loop inconsistency is high. We hypothesise that multiple G-loops, operating at different temporal scales and modalities, co-occupy a shared near-critical Ψ -band enabled by the brain’s hierarchical, modular architecture. Within this band, χ indexes cross-loop inconsistency/volatility and η indexes cross-loop competence and recovery ($F \rightarrow F^*$).

This proposal aligns with evidence for extended critical-like regions (Griffiths phases) in hierarchical brain networks, where power-law avalanches and large dynamic range persist over a band rather than a single point. The framework maps to known neural correlates, yields falsifiable predictions and a minimal simulation programme, and supports translational applications in pedagogy, cognitive therapeutics and AI.

SOC-for-Inference Hypothesis. *Homeostatic plasticity in hierarchical cortical networks drives neural dynamics into an extended near-critical Ψ -band because that regime maximises information capacity/dynamic range needed to minimise expected free energy during behaviour. Each G-loop tracks a difficulty set-point F^*F^* (thermostat) and selects policies by minimising GG; maintaining a non-zero target (tunable) F^*F^* preserves exploration/learning while SOC supplies the dynamical substrate that makes GG-minimisation effective.* ([PMC](#), [Nature](#))

Yes—framed per-loop, that’s empirically reasonable.

- **Local E:I balance is a core property of cortical and subcortical circuits.** Cortex is built from excitatory (glutamatergic) neurons and diverse inhibitory interneurons that **gate and sculpt** population dynamics; E:I is actively homeostated and varies by area/state—so each functional loop can have its **own operating E:I point**. ([PMC](#), [Frontiers](#))
- **E:I shifts change computation and are measurable.** Perturbing E:I alters information flow and oscillations; macroscopic proxies (e.g., LFP/EEG power–law slope) can track E:I changes in vivo, supporting loop-specific estimates. ([Institute for Neural Computation](#))
- **E:I tunes proximity to criticality.** Critical-regime benefits (dynamic range, information transmission/capacity) and the existence of an **extended critical band** in hierarchical brain networks imply that **E:I acts as a proximal knob** keeping each loop near critical—consistent with your Ψ -band account. ([SAGE Journals](#), [Nature](#))
- **Loops differ in their biophysics and neuromodulation.** Parallel cortico-basal ganglia–thalamo-cortical loops (motor/associative/limbic) have distinct receptor/topology profiles (e.g., D1/D2 gating of direct/indirect pathways), so the **precision-bias bb** and **E:I set-point** are **loop-specific**, not a single global constant. ([PMC](#))

So, modelling **each G-loop with its own E:I ratio** that interacts with your bias controller bb and temperature $T(\chi)T(\chi)$ is consistent with current evidence; the only caution is to

treat the “neutral” point and bandwidth as **context- and loop-dependent**, rather than universal. ([PMC](#), [Nature](#))

Each **G-loop** has its own *precision-bias* setting (your $b \approx \log D1:D2$ axis) that tilts processing toward **Control/Exploit** (D1-leaning, stable/compressive) or **Creative/Explore** (D2-leaning, flexible/decompressive). My earlier caution was only against assuming a **single, global “neutral” D1:D2 value** that applies to the whole brain across all tasks.

Why the per-loop reading is empirically reasonable:

- **Parallel cortico–basal-ganglia–thalamo-cortical loops** (motor, associative, limbic) are anatomically and functionally segregated; gating and dopamine effects are **loop-specific**, not global. ([PubMed](#), [Nature](#))
- **D1 vs D2** supports a **stability–flexibility** trade-off (dual-state theory), matching your exploit vs explore forks; the **optimal/neutral level depends on baseline and region** (the classic inverted-U and context dependence). So the “neutral” $b \approx 0$ will vary **by loop and state**. ([PubMed](#), [PMC](#))
- **Dopamine signalling is heterogeneous** across striatal/cortical territories (receptor distributions, kinetics, and circuit roles differ), reinforcing **loop-specific biases**. ([PMC](#))
- Computational and experimental work on BG **Go/NoGo** mechanisms shows DA-dependent gating that can differ **per circuit/task**, again arguing for b_k , not one brain-wide value. ([PubMed](#), [PMC](#))

So your diagram’s logic—**for each loop k** , branch by the pair $((r_k = F_k - F^*_k, b_k \approx \log \mathrm{D1:D2}))^{**}$ —is consistent with current evidence. If helpful, annotate the caption:

“**b-axis** = loop-specific precision bias ($\approx \log D1:D2$). **Neutral $b_k \approx 0$** is **context-dependent** (baseline DA, task, region). Exploit if $b_k > 0$; explore if $b_k < 0$; monitoring coordinates re-entry to Ψ when loops disagree.”

Why this overall model is defensible:

- **Active inference / EFE**. GG decomposes into predicted **risk**, **ambiguity**, and **epistemic value**; it is the policy-selection objective in contemporary Bayesian

accounts of action and perception. ([PMC](#), chrismathys.com)

- **Near-criticality benefits information processing.** Empirical and theoretical work shows **optimal information transmission, storage and dynamic range** near criticality; in hierarchical networks the critical “point” stretches to a **band** (Griffiths phase). ([PMC](#), [Nature](#), [PubMed](#))
- **Homeostasis → criticality.** In vivo and modelling studies indicate that **homeostatic rules tune cortical circuits toward critical dynamics**, consistent with a control process that keeps operation inside a band. ([PMC](#), [Nature](#))
- **E:I balance as the knob.** Adjusting excitation–inhibition moves the system toward/away from criticality and modulates encoding—exactly the sort of low-level control your loops would exploit. ([PubMed](#), [Physical Review Link](#))
- **Salience as coordinator.** The AI/dACC “salience network” causally coordinates switching between executive and default-mode systems, providing the monitoring/coordinating role at the branch. ([PMC](#), [PNAS](#))

How to write it

*We hypothesise that cortical SOC **evolved to support** active inference: homeostatic plasticity in hierarchical networks maintains a near-critical Ψ -**band** that maximises the information processing required to minimise expected free energy. Each G-loop uses a tunable difficulty set-point F^*F^* and minimises GG over policies; SOC supplies the dynamical regime that makes this feasible across scales.* ([PMC](#), [Nature](#))

What would count as evidence (and falsification)

- **Prediction 1:** Closer-to-critical telemetry (avalanche exponents/dynamic range) should correlate with **lower empirical proxies for GG** (e.g., lower prediction error + maintained model complexity) **at matched task difficulty**; pushing E:I away from balance should increase those proxies and impair performance. **Falsify** if moving toward criticality does **not** improve GG-proxies or task performance. ([PMC](#), [PubMed](#))
- **Prediction 2:** After perturbations, homeostatic adjustments should **restore** near-critical metrics **and** behavioural policies that lower GG; recovery speed indexes η . **Falsify** if homeostatic recovery does not return the system to the near-

critical band. ([PMC](#))

- **Prediction 3:** Across loops/scales, a **shared Ψ -band** should be observable (Griffiths-like), with salience-triggered transient “monitor” episodes during high cross-loop inconsistency χ that retune precision/temperature and re-enter the band. **Falsify** if multi-scale operation shows no common band or if salience activity does not precede re-entry. ([Nature](#), [PMC](#))

One-line guidance for the manuscript

“We treat F^*F^* as a **difficulty set-point** (thermostat) and GG as the **policy objective**; our claim is that **SOC provides the evolved dynamical regime** that maximises the information processing needed for GG-minimisation—and that homeostatic plasticity keeps multiple G-loops inside a **global Ψ -band**.” ([PMC](#))

strong, testable way to frame individual differences. In fact it gives you a clean “ **Ψ -dynamic-range hypothesis of g**”:

- **Cognitive Resilience axis (r):** individuals differ in the **range of challenge** $(F-F^*)(F-F^*)$ over which they can stay in-band (perform, learn, and recover) and how **quickly they can move the set-point F^*F^*** upward after success (η -driven thermostat).
- **Inference axis (b/T):** individuals differ in how **widely and precisely** they can modulate **precision-bias bb** ($D1 \leftrightarrow D2$) and **temperature $T(\chi)T(\chi)$** **without saturating**—i.e., how far they can shift between exploit and explore *and* still re-enter the Ψ -band promptly.

Together, those ranges form a person’s **Ψ -capacity** (general adaptive intelligence), which can **expand** with learning ($G_f \rightarrow G_c$ pipeline) as **η rises** and **F^*F^* ratchets** upward.

Individual-difference metrics (two-axis “ Ψ -profile”)

Let loop kk be in-band when $|\Delta F^k| \leq \epsilon_k$ $|\widehat{\Delta F}_k| \leq \epsilon_k$ and control variables aren’t saturated.

1. Resilience dynamic range

$R_r(k) = \max_{\Delta} \Delta(F_k - F_k^*)$ s.t. loop k stays in Ψ_k with accuracy/learning $\geq \theta$. $R^{\{(k)\}}_{\text{r}} = \max \Delta(F_k - F_k^*)$ s.t. loop k stays in Ψ_k with accuracy/learning $\geq \theta$.

Interpretation: how wide a *challenge window* a person can handle before disintegration.

2. Inference dynamic range

$R_i(k) = \max_{\Delta} \Delta(b_k, T_k)$ s.t. Ψ_k is re-entered within τ_{re} . $R^{\{(k)\}}_{\text{i}} = \max \Delta(b_k, T_k)$ s.t. Ψ_k is re-entered within τ_{re} .

Interpretation: controllable spread from **Control** to **Creative** (via b_b , T_T) without getting stuck.

3. Set-point mobility (thermostat gain)

$\mu F^*(k) = dF_k/d\eta_k$ (at matched volatility). $\mu^{\{(k)\}}_{F^*} = \frac{dF_k^*}{d\eta_k}$ (at matched volatility).

Higher μ_{F^*} = faster ZPD ratcheting as competence grows.

4. Valence-surprise span

Range of tolerable **positive/negative** prediction error (better-than/worse-than expected) before leaving Ψ_k . Captures “psychological buffer” for wins and setbacks.

5. Recovery half-life

$\tau_{1/2}(k)$: time to return $F_k \rightarrow F_k^*$ after a perturbation. Shorter = more resilient.

6. Band occupancy & hysteresis

Ω = time in Ψ_{global} $\Omega = \frac{\text{time in } \Psi_{\text{global}}}{\text{total}}$, and **hysteresis area** between Control \leftrightarrow Creative switching curves. Smaller area = cleaner meta-control.

Aggregate across active loops to get a **global Ψ -profile**:

$\Psi\text{-profile} = (R_r^-, R_i^-, \mu F^{*+}, \text{span}^-, \tau_{12}^-, \Omega, \text{hysteresis})$. $\mathbf{\Psi\text{-profile}} = \Big(\overline{R_{\text{r}}}, \overline{R_{\text{i}}}, \overline{\mu F^{*+}}, \overline{\text{span}}, \overline{\tau_{\frac{1}{2}}}, \overline{\Omega}, \text{hysteresis}\Big)$.

Prediction: higher **general adaptive intelligence** \approx larger $\overline{R_{\text{r}}}$, $\overline{R_{\text{i}}}$, $\overline{\mu F^{*+}}$, broader valence span, shorter recovery, higher occupancy, lower hysteresis—and these expand with targeted training (especially η -raising curricula that alternate Creative→Control passes).

One-paragraph insert

Ψ -dynamic-range hypothesis of g. We propose that individual differences in general adaptive intelligence reflect the **dynamic range** of the **Cognitive Resilience** axis (tolerable challenge window and η -driven mobility of $F \rightarrow F^{*+}$) and the **Inference** axis (controllable spread of χ and $T(\chi)$ without saturation), aggregated across loops in a global Ψ -band. As competence η rises through learning, $F \rightarrow F^{*+}$ **ratchets upward** ($G_f \rightarrow G_c$ consolidation), expanding both the subcritical “what we can automate” and the super/near-critical “how intense a challenge we can stabilise” ranges. We predict that people with larger two-axis ranges show higher **band occupancy**, **faster recovery** to $F \rightarrow F^{*+}$, **cleaner Control↔Creative switching** (low hysteresis), and **broader tolerance** to positive/negative surprise—and that these capacities are **trainable**.

Two-sentence add to your abstract (concise)

We further hypothesise that **individual differences in general adaptive intelligence** correspond to a person’s **Ψ -dynamic range**: the controllable span of (i) **challenge tolerance** on the Resilience axis and (ii) **precision-bias/temperature modulation** on the Inference axis, with **η -driven mobility of $F \rightarrow F^{*+}$** expanding this range via a $G_f \rightarrow G_c$ pipeline. This yields testable metrics (band occupancy, recovery half-life, hysteresis, set-point mobility) and concrete training targets to **expand** the Ψ -range.

What is well supported

1) Brains operate near critical regimes (a Ψ -like “near-critical band”), with signs of Griffiths-phase-like behaviour.

Neuronal avalanches and scale-free activity appear across species and modalities; network models and empirical data point to benefits for information capacity and dynamic range when networks sit near criticality. Evidence also suggests extended critical-like regimes (“Griffiths phases”) in heterogeneous networks, compatible with a *band* rather than a single point. ([PMC](#), [Nature](#), [Physical Review Link](#))

2) Excitation–inhibition (E:I) balance tunes proximity to criticality.

Perturbing E:I shifts avalanche statistics and oscillatory regimes, while balanced E:I supports near-critical dynamics. Human and animal work plus modelling converge on this principle. ([Physical Review Link](#), [PMC](#))

3) The salience network (AI/ACC) arbitrates network switching between executive/fronto-parietal control and default-mode systems.

Right anterior insula/ACC activity causally precedes switches, consistent with your “monitoring/arbitrator prong”. ([PMC](#))

4) Fluid intelligence (Gf) loads on the multiple-demand/fronto-parietal network; semantic/episodic knowledge (Gc-like functions) load more on DMN-centred systems.

Lesion, fMRI and network-neuroscience work link Gf to the MD/FPN system, while DMN contributes to memory/semantic representations often associated with Gc. ([Cell](#), [ScienceDirect](#), [Journal of Neuroscience](#))

5) Metastability and flexible network reconfiguration relate to cognitive flexibility and creativity.

The brain shows metastable dynamics at rest and during tasks; greater switching and dynamic reconfiguration are tied to creative ability. ([Nature](#), [PMC](#))

6) Dopamine D1 vs D2 supports a stability–flexibility trade-off (your exploit vs explore bias).

Prefrontal D1 favours robust, focused representations (stability/exploitation); D2 promotes flexibility/update. This “dual-state” view is well reviewed. ([PMC](#))

7) Arousal/uncertainty signals (LC–noradrenaline; pupil) modulate exploration–exploitation, consistent with a temperature-like controller $T(x)T(x)$.

Adaptive-gain theory and human pupillometry show that LC/NE state and pupil size track shifts between exploit and explore and respond to uncertainty/entropy—exactly the role you assign to $T(x)T(x)$. ([PubMed](#), [Nature](#))

8) Hierarchical predictive processing on multiple time-scales (many “loops” at different speeds) is a standard account.

Cortex appears organised along temporal hierarchies; ACC tracks volatility and adjusts learning—good support for your cross-timescale monitoring idea. ([PLOS](#), [PubMed](#))

9) “Compression vs. decompression” has plausible neural analogues.

Predictive/efficient coding frameworks formalise compression; hippocampal pattern separation vs. completion and successor-representation mapping provide concrete mechanisms for expanding vs. compressing structure (and for “map-based” inference). ([annualreviews.org](#), [PMC](#), [gershmanlab.com](#))

10) Boredom can drive exploration (a “soft” route into D2-like decompression).

Experimental and theoretical work suggests boredom functions as a regulatory signal pushing organisms to seek novelty/information, consistent with your “under-challenge” pathway. ([MDPI](#), [PubMed](#))

What is partly supported / still speculative

A) “Critical point = expected free energy (F)”.*

Expected free energy (often GG) is a decision-theoretic quantity in active inference that combines predicted risk/ambiguity to guide policy selection; *criticality* is a dynamical-systems property of neural population activity. There is intriguing theory linking neuromodulatory precision and EFE-guided behaviour, and suggestive dopaminergic findings, but equating the dynamical *critical point* with an EFE *set-point* is a proposal rather than an established empirical identity.. ([eLife](#))

B) A global “neutral” D1:D2 ratio as a formal bifurcation parameter across tasks.

There is strong support for $D1 \leftrightarrow \text{stability}$ and $D2 \leftrightarrow \text{flexibility}$ in *prefrontal circuits*. Each G-loop has its own precision-bias setting (your $b \approx \log D1:D2$ axis) that tilts processing toward Control/Exploit (D1-leaning, stable/compressive) or Creative/Explore (D2-leaning, flexible/decompressive). ([PMC](#))

C) χ as “inconsistency/conflict between loops” and η as “competence/return-to-setpoint” across loops.

There is solid evidence that ACC/insula/salience systems track conflict, volatility and uncertainty, and that learning rates and arousal adjust accordingly; competence-based intrinsic motivation is also supported. But the *specific* decomposition χ = cross-loop inconsistency and η = cross-loop return-to-critical-band is a promising synthesis rather than a direct empirical construct—worth operationalising in your telemetry. ([PubMed](#))

D) “Two D2 decompression routes”: boredom-driven vs error-driven.

Each route has independent support (boredom→exploration; volatility/surprise→exploration with increased arousal and ACC engagement), but showing both *within the same unified loop* is a prediction to test. ([MDPI](#), [PubMed](#))

Notes

- **Near-critical band:** “Large-scale brain activity exhibits signatures of near-critical dynamics, likely supported by balanced excitation–inhibition; heterogeneous network structure may create an extended critical-like band.” ([PMC](#), [Physical Review Link](#))
- **Trident arbitration:** “The salience network monitors surprise/uncertainty and arbitrates switches between default-mode (knowledge-rich) and fronto-parietal (control) modes.” ([PMC](#))
- **Exploit/explore biases:** “D1-dominant states stabilise representations; D2-dominant states facilitate updating / exploring—mapping naturally to exploit vs explore biases.” ([PMC](#))
- **Temperature/uncertainty:** “LC–NE arousal and ACC volatility signals adjust exploration in a temperature-like manner.” ([PubMed](#))
- **Maps and (de)compression:** “Hippocampal ‘predictive maps’ (successor representation) and pattern separation/completion implement the map update (decompress/recompress) you describe.” ([gershmanlab.com](#), [PMC](#))

What to treat as predictions (good targets for your Ψ -band telemetry)

- **Telemetry signatures of banded near-criticality** during skilled performance vs. learning (avalanche scaling, metastability maxima, E:I-sensitive spectral metrics). ([PMC](#), [Nature](#))
- **Prong-specific network states:** SN-led switches into FPN-dominated (exploit) vs DMN-coupled (creative/explore) episodes, with dopaminergic state (D1/D2-leaning tasks or pharmacology) biasing the branch. ([PMC](#))
- **Two exploration routes:** (i) boredom/under-challenge → moderate arousal, increased switching; (ii) error/volatility → higher arousal, ACC-linked learning-rate/temperature upshift. ([MDPI](#), [PubMed](#))
- **Map (de)compression:** hippocampal SR-like updates and separation/completion patterns tracking your Phase-3 decompression/recompression claims. ([gershmanlab.com](#), [PMC](#))

Think in two layers:

1) A global Ψ -band (system level)

The brain's hierarchical, heterogeneous architecture supports an **extended near-critical regime** in which many subsystems can co-exist near criticality together. That's your Ψ_{global} .

2) Loop-specific Ψ -bands (per G-loop)

Each active loop k has its own admissible near-critical window, Ψ_k , around its set-point F_k^* . A practical definition:

$$\Psi_k = \{ |\Delta F_k| \leq \varepsilon_k, \chi_k \in [\chi_k^{\text{low}}, \chi_k^{\text{high}}], |b_k| < b_k^{\text{max}}, T(\chi_k) \text{ not saturated} \} \Psi_k$$

$$\varepsilon_k = \frac{1}{\text{Big}} \left(\frac{1}{|\widehat{\Delta F}_k|} \right) \leq$$

$$\varepsilon_k, \chi_k \in [\chi_k^{\text{low}}, \chi_k^{\text{high}}], |b_k| < b_k^{\text{max}}, T(\chi_k) \text{ not saturated} \} \Psi_k$$

Optionally add telemetry constraints (pick what you'll measure):

- avalanche/branching near critical (e.g., $\sigma_k \approx 1$; power-law exponent in a target interval),
- peak metastability/dynamic range for that loop's network,
- stable map compression metrics (η up; no runaway complexity).

3) How the two layers relate

- **Global Ψ** is the **intersection** where all *currently engaged* loops sit within their Ψ_k and **agree enough** across scales: high cross-loop consistency ρ (i.e., low meta-uncertainty χ_{meta}).
- When loops **conflict** ($\chi_{\text{meta}} \uparrow$), the **Monitoring/Salience** prong acts: brief orient/reset, adjust λ , centre b , reduce λ (quarantine propagation), or nudge one loop's F_k^* so the set of loops re-enters Ψ_{global} .
- Coupling means each Ψ_k is **context-dependent**: neighbouring loops (and arousal/neuromodulators) can shrink/shift its window.

4) Why this is empirically reasonable

- Near-critical **bands** (not just points) arise naturally in heterogeneous, hierarchical networks; modules can be locally closer/further from criticality yet remain inside a

shared extended regime.

- Different scales/modalities (your multiple G-loops) mapping to partially distinct networks is exactly where such local bands would live.
- Your meta-parameters fit:
 - χ = cross-loop inconsistency/volatility (drives TT and monitoring actions),
 - η = competence/recovery (how quickly $F_k \rightarrow F_k * F_k \rightarrow F_k^*$ while maintaining useful compression).

5) How to write it in the paper (one sentence)

“We posit a **global Ψ -band** within which multiple G-loops co-operate, each with its own **loop-specific Ψ_k** around $F_k * F_k^*$; χ tracks cross-loop inconsistency and η indexes recovery of each loop toward its set-point, allowing the Monitoring/Salience prong to keep the ensemble within Ψ -band operation.”

That keeps the global picture and gives you concrete, testable per-loop criteria.

As a **working model**, treating each cortical column (or local microcircuit) as a G-loop and the interactions **between** columns as an **extended Ψ -band/Griffiths-phase zone** is empirically reasonable. There's direct evidence for critical-like dynamics at **column/layer scale**, and strong evidence that **hierarchical, modular networks** (like cortex) stretch a single critical point into a **band**, which is exactly the substrate your χ/η meta-parameters would sit on. The main caveat is that “columns” are clearer in some areas (e.g., barrel cortex) than others—so when you generalise beyond primary sensory cortex, say “local microcircuits/modules” rather than assuming a universal, uniform column. ([Frontiers](#), [PMC](#))

Why this mapping is defensible

- **Critical-like dynamics inside a single column/layer.** Neuronal avalanches have been observed across cortical layers and, notably, **within a single rat barrel column** in vivo; classic slice/awake studies likewise show avalanche statistics at mesoscopic scales. That supports a **per-loop (per-column) near-critical regime**. ([Frontiers](#), [PubMed](#), [PMC](#))
- **E:I balance tunes proximity to criticality.** Balanced excitation–inhibition makes avalanches and oscillations co-emerge at a critical state; pushing E:I off-balance moves the system away from criticality—consistent with a **loop-specific control knob**. ([Journal](#))

[of Neuroscience](#))

- **From columns to a shared band (Griffiths phase).** In **hierarchical, modular networks**, criticality is **stretched into a band** (a Griffiths phase) rather than a knife-edge point; this has been demonstrated on synthetic hierarchies **and** empirical connectomes. That supports your **global Ψ -band across many local loops**. ([Nature](#), [materias.df.uba.ar](#))
- **Columns/microcircuits are real—but heterogeneous.** Modern reviews describe a canonical microcircuit “theme with variations” across areas/species. Others caution that a **single** column concept is not a universal functional unit. So: model **per-loop** where columnar structure is strong (e.g., barrels, A1), and talk about **local modules** elsewhere. ([PMC](#))

Where χ and η fit

Your proposal— χ as cross-loop inconsistency/volatility and η as cross-loop competence/recovery to each loop’s set-point $F_k \rightarrow F_k^*$ —is not a standard empirical metric, but it’s **compatible** with: (i) hierarchical predictive processing on multiple time-scales and (ii) salience/ACC roles in monitoring volatility and reconfiguring networks. It’s a **good, testable** modelling layer on top of the established criticality picture. ([PMC](#))

Concrete tests (column \leftrightarrow band)

1. **Column-level criticality + global band:** show avalanche exponents/branching ratios **within single columns** sit in a **common interval** while the animal performs a task; perturb E:I locally and test that the column leaves the band and that χ (cross-column inconsistency) rises until re-entry. ([Frontiers](#), [Journal of Neuroscience](#))
2. **Dynamic-range prediction:** nearer-critical columns should display **higher dynamic range** and information transmission; correlate these with performance and with your η (faster recovery of $F_k \rightarrow F_k^*$ to F_k^*). ([SAGE Journals](#), [PubMed](#))
3. **Griffiths-phase signature:** across columns/modules, look for **power-law regimes** and **rare-region effects** consistent with a **band** (varying exponents across modules) rather than a single critical point. ([Nature](#))

One critical distinction in the large scale networks is between the hippocampal-prefrontal value-landscaped cortical maps (which are needed for higher level relational reasoning) associated with long-term working memory in FPCN-A, and the FPCN-B fronto-parietal network for short-term (classic) working memory for rules application/rapid rule learning and the standard limited capacity working memory,

Where it lives in the loop

- **Phase 3 — Map (ϕ):** builds/updates a *predictive* hippocampal–prefrontal map (successor-representation style) and overlays value. This is the long-horizon, relational “**long-term working memory**” workspace. Network-wise, it leans on **hippocampus/DMN coupling with FPCN-A**, which is more internally/future-oriented and interfaces with memory systems. (gershmanlab.com, [PubMed](#), [eLife](#), [Journal of Neuroscience](#))
- **Phases 4–5 — Decide/Gate → Act:** hold/apply rules, update on the fly, and sequence actions. This is the classic **short-term working memory / rule application** workspace, dominant in **FPCN-B** with **DAN** support. FPCN-B sits nearer sensorimotor control, whereas FPCN-A couples more to DMN/long-term memory—dissociations shown across tasks. ([PNAS](#), [Journal of Neuroscience](#))
- **Monitoring/Saliency (AI/dACC):** arbitrates when to hand off between these workspaces (and briefly performs orient/reset/reconfigure). Causal evidence shows the saliency network triggering switches between default-mode and executive/control systems. ([PNAS](#), [PubMed](#))

How controllers steer the hand-off

- **Precision-bias $b \approx \log \frac{D1}{D2} \approx \log \frac{\text{D1}}{\text{D2}}$:** $b \uparrow$ (stability/exploit) biases toward **FPCN-B + DAN** (rule maintenance/closure); $b \downarrow$ (flexibility/explore) biases toward **FPCN-A + hippocampus/DMN** (relational recombination/planning). Empirically, FPCN-A shows stronger coupling with memory-oriented systems, FPCN-B with sensorimotor/attention control. ([PNAS](#), [Journal of Neuroscience](#))
- **Temperature $T(\chi)T(\chi)$:** higher uncertainty/volatility ($\chi \uparrow$) raises policy temperature, favouring exploratory sampling and DMN–FPCN-A engagement; lower T helps lock in FPCN-B-centred rule execution. ([eLife](#))

Why this bridges micro ↔ macro

Your “scale-free” picture holds: local (column/microcircuit) inference loops feed into these large-scale workspaces, and the same controllers ($b, T(\chi)$) plus saliency-driven arbitration organise behaviour across scales. Hippocampal SR-style maps provide the *relational substrate*;

FPCN-B provides the *rule/sequence executor*; salience ensures *the right workspace at the right time*. (gershmanlab.com, [PNAS](#))

Reasoning

Where each reasoning mode lives

Deduction → Control prong (Exploit/Compress)

- **When:** Phase 3 (**Map**) during a Control branch; again in **Phase 6 (Update)** for consistency checks.
- **What it does:** Rule-based closure on the current map ϕ to derive consequences, eliminate inconsistent options, and tighten beliefs.
- **Controllers:** $b \uparrow$ (stability bias), TT low–mid, $\lambda \uparrow$ (propagate validated structure).
- **Networks:** FPCN-B (+ DAN) for rule maintenance/sequence execution.
- **Operational hook:** apply a closure operator until a budget is met
 $\phi \leftarrow \text{DED}(\phi, R; \text{budget } \kappa, \text{strictness } \sigma)$
 $\phi \leftarrow \text{DED}(\phi, R; \text{budget } \kappa, \text{strictness } \sigma)$
(You can keep κ, σ implicit if you want to stay minimalist.)

Induction → Control → Update (Compress/Generalise)

- **When:** Phase 3 (**Map**) after successful deductions; **Phase 7 (Consolidate)**.
- **What it does:** Compress patterns, widen generalisation, raise competence η .
- **Controllers:** $b \uparrow$, $T \downarrow$, then $\lambda \uparrow$ if validation passes.
- **Networks:** Hippocampus ↔ FPCN-A (pattern abstraction) handing to FPCN-B for rule deployment.
- **Operational hook (MDL/predictive gain):**
 $\Delta \text{Compression} = \text{MDL}_{\text{old}} - \text{MDL}_{\text{new}} \Rightarrow \eta + = wD1 \Delta \text{Compression}$

$$\mathrm{MDL}_{\text{old}} - \mathrm{MDL}_{\text{new}} \quad \rightarrow \quad \eta_{\text{rel}+} = w_{D1} \Delta \text{Compression}.$$

Abduction → Creative prong (Explore/Decompress)

- **When:** Phase 3 (**Map**) during a Creative branch; also during **Monitoring** micro-loops when ambiguity is high.
- **What it does:** Proposes new hypotheses/structures that could explain the evidence (or the mismatch).
- **Controllers:** $b \leftrightarrow \downarrow b$, $T \uparrow T$, $\lambda \downarrow$ (sandbox), optional novelty budget $d \uparrow d$.
- **Networks:** Hippocampus/DMN + FPCN-A (relational recombination, analogy).
- **Operational hook (posterior-sparse search):**

$$H^* \in \arg \max_{H \in \mathcal{H}} H[\log P(D|H) + \log P(H)]$$
subject to a proposal budget d .

Counterfactuals → Decide & Act (Plan/Intervene/Test)

- **When: Phase 4 (Decide)** to choose *informative* actions; **Phase 5 (Act)** to run A/B probes; also in Phase 1 (Sense) as short rollouts.
- **What it does:** Simulates “what-if” outcomes under alternative policies or interventions to reduce ambiguity and disambiguate hypotheses.
- **Controllers:** use $T(\chi)T(\chi)$ to prioritise information gain when $\chi \uparrow$; keep χ centred if monitoring.
- **Networks:** Hippocampus (predictive map/SR) + FPCN-A for long-horizon simulation; FPCN-B executes the chosen test.
- **Operational hook (EFE with epistemic term):**

$$\pi^* \in \arg\min_{\pi} G(\pi) \approx \arg\max_{\pi} \text{EIG}(\pi) - \text{counterfactual value} - \text{risk} - \text{cost}$$

Absolutely—**analogy** is the bridge between Creative (abductive propose) and Control (deductive/inductive verify & compress). Here's **drop-in text** that matches your scheme and notation.

Analogy (Relational Mapping) → Creative→Control bridge

- **When:** Phase 3 (**Map**) during a **Creative** branch (retrieve & align); hands off to **Control** for constraint-checking and consolidation when a mapping passes tests; also recruited in **Monitoring** micro-loops when ambiguity is high.
 - **What it does:** Retrieves a *source* structure by **relational similarity**, aligns it to the *target* (structure-mapping), **projects** candidate inferences/policies, then **tests** them (deduction/counterfactuals) and **abstracts** a schema (induction) if validated.
 - **Controllers:** retrieval/alignment with $b \approx 0$ or $\downarrow b \approx 0$ or \downarrow , $T \uparrow$, $\lambda \downarrow$ (sandbox, optional $d \uparrow$); after validation, **commit** with $b \uparrow$, $T \downarrow$, $\lambda \uparrow$.
 - **Networks:** Hippocampus (pattern completion/separation; SR “predictive map”), **FPCN-A** + **DMN/mPFC** for relational mapping and recombination; **FPCN-B** (+ DAN) for constraint checking and execution/roll-out.
 - **Operational hook (structure-mapping + projection):**

$$M^* \in \arg \max_{M \in \mathcal{M}} [\text{Simrel}(M) + w_{\text{feat}} \text{Simfeat}(M) - \beta \text{Viol}(M)] \text{ s.t. proposal budget } d, M^* \in \mathcal{M}$$

$$\arg \max_{M \in \mathcal{M}} [\text{Sim}_{\text{rel}}(M) + w_{\text{feat}} \text{Sim}_{\text{feat}}(M) - \beta \text{Viol}(M)] \quad \text{s.t. proposal budget } d,$$

$$IT \leftarrow \text{Project}(IS, M^*), \text{score via EV or } -G, \text{ then verify (deduction) and, if successful, abstract (induction).}$$

$$I_T \leftarrow \text{Project}(I_S, M^*), \quad \text{score via EV or } -G, \text{ then verify (deduction) and, if successful, abstract (induction).}$$
-

Threading analogy through the 8 phases (add these clauses)

- **1) Sense/Move:** cue source retrieval by relational/contextual probes; run brief counterfactual roll-outs for top candidates.
- **2) Branch:** when $\Delta F = b(F - F^*) \geq 0$ and χ high, prefer **Creative**; Monitoring may call an *analogy probe* when χ_{meta} is high.
- **3) Map (Creative):** retrieve→align→project ($S \rightarrow T$) ($S \rightarrow T$) under $\lambda \downarrow$, $T \uparrow$, $b \approx 0$; **Map (Control):** verify→compress successful projections, raise η .

- **4) Decide:** pick the **most informative/low-risk** analogical projection using Softmax over EV (or $-G-G$) with $T(\chi)T(\chi)$ and bb .
- **6–7) Update→Consolidate:** credit η if the analogical policy **reduces future surprise**; **schema abstraction** widens $\lambda \uparrow$ for transfer.

How it threads through the 8 phases (at a glance)

- **1) Sense/Move** – quick counterfactual rollouts to score candidate policies; compute $F, \chi F, \chi$.
- **2) Branch** – if $\Delta F = b(F - F^*) \geq 0$ and χ low → **Control** (deduce/compress); if χ high → **Creative** (abduce/decompress); if χ_{meta} high → **Monitoring** (brief orient/reset, then re-branch).
- **3) Map** – **Control: deduction** → **induction** (close rules, compress, raise η); **Creative: abduction** (new H, sandboxed by $\lambda \downarrow, d \uparrow$).
- **4) Decide** – **counterfactual selection** via Softmax over EV (or $-G-G$) with $T(\chi)T(\chi)$ and bb .
- **5) Act** – run the test/intervention; gather evidence.
- **6) Update** – credit η for *either* successful compression (D1) *or* successful strategy change (D2) that reduces future surprise; adjust b, T, F^*, T, F^* .
- **7) Consolidate** – promote validated structure ($\lambda \uparrow$); keep speculative bits quarantined.
- **8) Reset** – restore near-critical windows; prep for next cycle.

Controller intuition (why this makes sense)

- bb steers **deduction/induction** (when $b \uparrow$) vs **abduction/counterfactual testing** (when $b \downarrow$ and $T \uparrow$).

- $T(x)T(\chi)$ ensures that when uncertainty/conflict rises, you *sample more broadly* (counterfactuals) and consider **abductive** alternatives.
- η grows from **both**: (i) D1 compressive wins *and* (ii) D2 successful escapes/re-commitments—your lived-experience point.
- F^*F^* sets challenge so you don't collapse into trivial deduction *or* endless abduction; the loop hovers in the Ψ -band.

One-sentence drop-in

*Reasoning is phase-specific in the G-Loop: **deduction** and **induction** dominate the Control path (compress/propagate), **abduction** is the Creative path's engine (decompress/propose), and **counterfactual simulation** selects informative actions at Decide/Act; Monitoring coordinates brief orient–reset–reconfigure episodes when cross-loop inconsistency is high. Analogy operates as a **Creative→Control bridge**: retrieve and align a relationally similar source, project counterfactual inferences, then verify and compress validated mappings into the rule set.*

Constraint satisfaction (feasibility & propagation)

When:

- **Phase 3 (Map)** on the **Control** branch (primary locus), and as a **feasibility gate** on the **Creative** branch before proposals leave the sandbox.
- **Phase 4 (Decide)** as a **feasible-set filter** on policies.
- **Monitoring** triggers constraint checks when conflict/high risk is detected.

What it does:

- Enforces **hard invariants** (must hold) and **soft preferences** (penalised if violated).
- **Propagates** constraints to prune the search space (like CSP/arc consistency), tightening beliefs and reducing ambiguity.
- Provides a **safety envelope** for Creative moves and counterfactual tests.

Controllers:

- **Control pass:** $b \uparrow$, $T \downarrow$, $\lambda \uparrow$ (propagate validated constraints widely).
- **Creative pass (sandbox):** $b \approx 0$ or $b \downarrow$, $T \uparrow$, $\lambda \downarrow$ (propose freely **inside** a safety envelope).
- **Monitoring:** briefly recentres b , adjusts $T(\chi)$, and tightens λ if a violation is imminent.

Networks:

- **FPCN-B + dACC/AI (Salience)** for rule enforcement, error-likelihood and conflict monitoring;
- **Hippocampus/FPCN-A/DMN** when constraints are **relational** (analogy/schema) and must be aligned before testing.

Operational hooks (minimal maths)

Let **constraints** be functions on hypotheses H and policies π :

- **Hard:** $g_j(H) \leq 0$, $g_j(\pi) \leq 0$ (invariants/safety).
- **Soft:** penalties $c_j(H) \geq 0$, $c_j(\pi) \geq 0$ (preferences/limits).
- **Chance constraints (optional):** $P[g_j(\pi) \leq 0] \geq 1 - \epsilon$ (for all j).

Feasible sets

$H_{\epsilon} = \{H: P[g_j(H) \leq 0] \geq 1 - \epsilon \forall j\}$, $\Pi_{\epsilon} = \{\pi: P[g_j(\pi) \leq 0] \geq 1 - \epsilon \forall j\}$.
 $\mathcal{H}_{\epsilon} = \{H: P[g_j(H) \leq 0] \geq 1 - \epsilon \forall j\}$, $\mathcal{\Pi}_{\epsilon} = \{\pi: P[g_j(\pi) \leq 0] \geq 1 - \epsilon \forall j\}$.

Map-level constraint propagation (Control)

$(\phi, H_{\epsilon}) \leftarrow \text{CONS}(\phi, H_{\epsilon}; \text{budget } \kappa, \text{strictness } \sigma)$,
 $\mathcal{H}_{\epsilon} \leftarrow \text{CONS}(\phi, \mathcal{H}_{\epsilon}; \text{budget } \kappa, \text{strictness } \sigma)$.

which narrows the hypothesis space and tightens beliefs; raise η when propagation **reduces future surprise/ambiguity** at the same $F \cdot F^*$.

Creative proposals with feasibility gate

$H' \sim \text{Propose}(\phi; d, \lambda \downarrow, T \uparrow)$, accept only if $H' \in H_{\leq \epsilon}$ (or minimal $\text{Viol}(H')$). $H' \sim \text{Propose}(\phi; d, \lambda \downarrow, T \uparrow) \quad \text{accept only if } H' \in \mathcal{H}_{\leq \epsilon} \text{ (or minimal Viol}(H'))$.

Decide with constrained objective (soft constraints folded into the score)

$G \sim (\pi) = G(\pi) + \sum_j \mu_j c_j(\pi) \Rightarrow \pi^* \in \arg \min_{\pi \in \Pi} G \sim (\pi)$, $\tilde{G}(\pi) = G(\pi) + \sum_j \mu_j c_j(\pi)$
 $\quad \rightarrow \quad \pi^* \in \arg \min_{\pi \in \Pi_{\leq \epsilon}} \tilde{G}(\pi)$,

then select by Softmax with $T(\chi)T(\chi)$ and bias bb . (Use μ_j as multipliers to avoid conflict with your λ generalisation width.)

Safety envelope (action guard)

Project proposed actions/policies onto the feasible set before execution:

$\pi \leftarrow \Pi_{\leq \epsilon}[\pi]$ (nearest feasible “trust region”). $\pi \leftarrow \Pi_{\leq \epsilon}[\pi]$
 $\Pi_{\leq \epsilon}[\pi] = \arg \min_{\pi' \in \Pi} \|\pi - \pi'\| \quad \text{subject to } \pi' \in \Pi_{\leq \epsilon}$

How it threads through the 8 phases (add-on lines)

- **1) Sense/Move:** sample observations targeted at **active constraints** (query the variables that disambiguate feasibility).
- **2) Branch:** high χ and near-zero $\Delta \hat{F}$ may invoke **Monitoring** to run fast feasibility checks before allowing a Creative pass.
- **3) Map:**
 - **Control:** run **constraint propagation**; if consistent, **deduction** \rightarrow **induction**; $\lambda \uparrow$ once validated $\rightarrow \eta \uparrow$.
 - **Creative:** **propose** \rightarrow **feasibility filter** \rightarrow **test**; keep $\lambda \downarrow$ until constraints pass.
- **4) Decide:** **constrained Softmax** over $\Pi_{\leq \epsilon}$; soft penalties in \tilde{G} .
- **5) Act:** enforce a **safety guard**; abort or re-plan if a predicted hard constraint is at risk.

- **6) Update:** strengthen learned constraints that repeatedly validate; relax or refine soft ones that block performance without reducing future surprise.
- **7) Consolidate:** propagate **stable constraints** broadly ($\lambda \uparrow \lambda \text{lambda} \uparrow$); keep uncertain ones local.
- **8) Reset:** retain invariants; clear temporary/task-specific soft constraints unless promoted.

One-sentence insert (for your summary line)

*Constraint satisfaction is implemented as **feasibility gating and propagation**: hard/soft constraints define feasible hypothesis/policy sets; Creative proposals are sandboxed by a feasibility filter, and Control performs constraint propagation and consistency checks before induction and global propagation.*

This way, your everyday “does it satisfy x, y, z?” reasoning is explicitly part of the **Control prong**, **guards** the **Creative** prong, and cleanly integrates into **Decide/Act**—all with the same χ , $T(\chi)$, λ , F^* , χ , and η machinery you already have.

Why reasoning helps (in G-Loop terms)

- **Abduction (Creative prong) → hypothesis generation under uncertainty.**
When $\chi \uparrow$ and $\Delta F \geq 0$, abductive moves let the loop **escape local optima** and propose explanations that can *reduce future GG* after testing. Evolutionary payoff: rapid **strategy change** in volatile environments; the seed of innovation.
- **Counterfactual simulation (Decide→Act) → safe, information-efficient exploration.**
Running “what-ifs” before acting gives **epistemic value** without costly errors. Payoff: higher **sample-efficiency** and **risk control** (hunt, tool use, social manoeuvring) while keeping the loop near the Ψ -band.
- **Deduction (Control prong) → reliable closure and coordination.**
Once a candidate structure works, deductive closure **compresses and stabilises** it (D1), enabling **precise, multi-step execution** (planning, syntax, techne). Payoff: lower

on-line cost, reliable cooperation, cumulative routines.

- **Induction (Control→Consolidate) → compact, transferrable knowledge.**
Pattern abstraction (MDL/predictive gain) **raises η** and **widens λ** , turning specific wins into **generalisable skills**. Payoff: **far transfer** and cumulative culture.

Where working memory comes in

- **“Long-term working memory” (Phase 3 Map; FPCN-A ↔ hippocampus/DMN).**
Holds **value-landscaped relational maps** and supports analogy, abstraction, multi-step **counterfactual rollouts** (long horizon). This is the big human upgrade to Phase 3.
- **“Short-term working memory” (Phases 4–5; FPCN-B + DAN).**
Maintains **rules/goals** and sequences actions during execution, enforcing closure and rapid rule learning. This is the executor that **locks in** compressive wins.
- **Monitoring/Salience (AI/dACC).**
Detects cross-loop inconsistency ($\chi_{\text{meta}}\chi_{\text{meta}}$), **re-centres \mathbf{bb}** , adjusts $T(\mathbf{x})T(\chi)$, and **hands off** between the two WM workspaces at the branch. This keeps the whole system inside the **Ψ -band** rather than tipping into rigidity or chaos.

The evolutionary punchline (one paragraph you can paste)

Evolutionary rationale. We propose that human reasoning is an **expansion of the G-Loop’s Map→Decide core** that improves survival by **minimising expected free energy** more efficiently in complex, changing niches. Abduction and counterfactual simulation (Creative path and Decide) **extend search** in hypothesis and policy space while controlling risk; deduction and induction (Control path and Consolidate) **compress and propagate** validated structure for reliable execution and cultural transmission. Two coordinated working-memory systems implement this division of labour: a **relational, value-landscaped map** (hippocampus with FPCN-A) for long-horizon reasoning, and a **rule/sequence buffer** (FPCN-B) for short-term maintenance and action. The salience system arbitrates the hand-off, keeping loops within a **near-critical Ψ -band** where information capacity and dynamic range are high. On this view, human reasoning did not replace the ancestral loop; it **deepened Phase 3 and sharpened Phases 4–5**, yielding more competent η gains from both **compression** and **successful strategy change**.

Which aspects of reasoning “fit” the evolutionary benefit?

- **Counterfactual depth:** increases informative action selection (lower GG per sample).
- **Relational compositionality:** supports **far transfer** (bigger $\lambda \uparrow$ when validated).
- **Flexible meta-control:** fast **b/T** retuning avoids lock-in and catastrophic exploration.
- **Cumulative culture:** deduction/induction make skills **stable, shareable programmes** (high $\eta \rightarrow$ robust routines).
- **Innovation under safety:** abduction + counterfactuals generate novelty **without** excessive real-world cost.

Quick prediction hooks

- **Mode-specific WM coupling:** Creative episodes show **FPCN-A ↔ hippocampus/DMN** coupling; Control episodes show **FPCN-B ↔ DAN** dominance; Monitoring precedes the switch.
- **Controller signatures:** abduction/counterfactual phases exhibit $T(x) \uparrow, T(\chi) \uparrow, b \downarrow \approx 0$; deduction/induction show $T \downarrow, b \uparrow$.
- **Competence growth:** η increases both after **compressive wins** and after **successful policy changes** that reduce future error at matched F^*F^* .

I’m taking “near-critical band” results and specifying a **Trident-style branching mechanism**: a **non-bifurcated subcritical shaft** (autopilot), a **branch point** near $F \approx F^*$, and **two super(near)-critical prongs—Control** (compress/exploit) and **Creative** (decompress/explore)—**coordinated by Salience/attention**. Outside the Ψ -band, the dynamics collapse into **over-synchronised rigidity** or **over-chaotic fragmentation**. That is indeed beyond standard models and it’s a crisp, testable extension.

Here’s clean, drop-in language and a lightweight formal cartoon to make it concrete:

Drop-in: Trident branching (conceptual & formal)

Concept.

- **Subcritical shaft (Autopilot):** for $F \approx F^* < 0$ and low χ , the loop sits in a **single, stable** regime—fast, fluent, consolidated (Gc-like) processing without a branch.
- **Branch point (near-critical):** as **precision-weighted gap** $\Delta F = b(F - F^*)$ rises to ≥ 0 and χ lifts, the system enters a **bifurcation zone** where two coordinated modes are available; **Salience** arbitrates and briefly executes orient–reset–reconfigure.
- **Super(near)-critical prongs:**
 - **Control (Exploit/Compress; D1-leaning):** partial synchrony, rule closure, induction, propagation ($\lambda \uparrow$).
 - **Creative (Explore/Decompress; D2-leaning):** metastable/“chimera-like” episodes, abduction/analogy, counterfactual testing ($\lambda \downarrow$, $T \uparrow$).
- **Band exit (failure modes):** outside Ψ the system **locks-in** (over-synchronised rigidity) or **fragments** (over-desynchronised chaos). We treat clinical phenotypes (e.g., compulsive rigidity; disorganised thought) as **hypothesis-level analogues**, not identities.

Normal-form caricature (for readers who like equations).

Use two low-dimensional coordinates:

- R = **integration/synchrony** order parameter (integration vs fragmentation),
- y = **mode coordinate** (Control vs Creative, $y > 0$ vs $y < 0$).
Let the **bifurcation controls** be $u = F - F^*$ (challenge gap) and $v = b$ (precision-bias). A minimal potential:

$$V(R, y) = \beta R^4 - \mu(u, \chi) R^2 + 14y^4 - u^2 y^2 - v y$$

$$V(R, y) = \frac{\beta}{4} R^4 - \frac{\mu(u, \chi)}{2} R^2 + \frac{1}{4} y^4 - \frac{u}{2} y^2 - v y$$

with stochastic gradient flows $\dot{R} = -\partial V / \partial R + \zeta_R$ and $\dot{y} = -\partial V / \partial y + \zeta_y$.

- For $u < 0$ (subcritical), y has **one** minimum (no branch): **shaft**.
- Near $u \approx 0$, a **cusp/pitchfork-like** region opens in y : two attractors (Control/Creative); $v = b$ tilts the branch.
- **Salience** acts as a **fast controller** that recentres $v \rightarrow 0$, briefly lifts **temperature** $T(\chi)$, or nudges μ to maintain **moderate RR**—keeping the trajectory **inside** the Ψ -band.
- For **large** $|u|$, either RR saturates high (rigidity) or collapses low (fragmentation): Ψ -band **exit**.

This normal form is only an **analytical cartoon**—your network-level instantiation (e.g., Kuramoto/Wilson–Cowan on a hierarchical-modular graph) provides the ground truth where:

- **Control** aligns with **higher, stable RR** and compressive updates,
- **Creative** aligns with **metastable RR** and decompressive search,
- **Monitoring** issues **short pulses** (centre bb , adjust $T(\chi)$, tighten λ) to re-enter Ψ after divergence.

How this extends prior work (and stays testable)

- Prior criticality work gives you the **band** and **benefits** (dynamic range, information capacity) but not the **branching logic**. The Trident adds the **decision-level geometry** and **controllers** that explain *when* and *how* the system chooses compress vs decompress—and *how* it avoids falling out of Ψ .
- It yields **quantitative predictions**: within-band you should see **metastability peaks**, **power-law episode durations**, and **SN-led** switching into FPN-B (Control) vs FPCN-A/DMN (Creative). Outside the band you should see **over-synchrony** (rigidity, narrow λ , high hysteresis) or **over-desynchrony** (unstable RR, poor re-entry).
- Clinical mappings remain **tentative**: hypothesise **locked-in RR** and persistently high bb (plus low TT) in compulsive rigidity; hypothesise **unstable RR**, high TT , and poorly centred bb in disorganisation. These are **testable computational phenotypes**, not diagnoses.

Two sentence insert (Discussion)

*We formalise a **Trident** branching in which a **single, non-bifurcated subcritical regime** (autopilot) meets a **near-critical branch point** governed by the precision-weighted gap $\Delta F^{\wedge}=b(F-F^*)\widehat{\Delta F}=b(F-F^{\wedge*})$ and uncertainty $\chi\backslash\chi$. Saliency/attention coordinators route dynamics into **Control** (compress) or **Creative** (decompress) prongs and deliver brief re-centring pulses; outside the Ψ -band the system either **over-synchronises** (rigidity) or **over-desynchronises** (fragmentation), providing clear, falsifiable failure modes.*

Yes—both fit, and they're complementary.

When to use which

- **Kuramoto (phase-only, macroscale):** Ideal for modelling the Ψ -band as partial synchrony/metastability on a **connectome or hierarchical-modular network**. It reproduces **smeared/smooth crossovers** with **power-law desynchronisation durations** (exponents in the $\sim 1\text{--}2$ range) on human connectomes and HMNs—i.e., **extended dynamical criticality** consistent with a near-critical **band**, not a knife-edge point. Map your controllers as: coupling/homeostasis $\leftrightarrow F^*F^{\wedge*}$ thermostat; precision-bias $bb \leftrightarrow$ subnetwork gain tilt (FPN-B vs DMN/FPCN-A); $T(\chi)T(\backslash\chi) \leftrightarrow$ noise/frequency jitter.
- **Wilson–Cowan (E/I, mesoscale):** Use when you want **explicit excitation–inhibition** and local microcircuit realism (columns/nodes each with E/I populations). You can sit each node near a **Hopf edge** (critical slowing/oscillations), control **E:I** directly, and recover phases from E-activity to feed cross-node coupling. Then the same G-Loop knobs apply: E:I/homeostatic gain $\leftrightarrow F^*F^{\wedge*}$; bias bb skews pathway gains; $T(\chi)T(\backslash\chi)$ scales noise/input variance. This gives you a mechanistic route from E/I to band occupancy while still supporting Kuramoto-like order-parameter readouts (e.g., global RR).

A neat hybrid (what I'd recommend)

- **Node dynamics:** Wilson–Cowan E/I per node (for local E:I and Hopf-edge control).
- **Network coupling:** Kuramoto-style phase coupling between nodes (cheap, captures **partial synchrony** and **metastability** cleanly).
- **Controllers:**

- $F \cdot F^*$ thermostat \rightarrow slow homeostatic adjustment of node gains/coupling toward a target $R \cdot R^*$ (partial synchrony);
- bb (D1:D2 tilt) \rightarrow multiplicative gain to Control vs Creative subnetworks;
- $T(\chi)T(\chi)$ \rightarrow exploration temperature (noise on inputs/frequencies) raised when cross-loop inconsistency χ is high.
- **Branch logic:** Near $\Delta F = b(F - F^*) \approx 0$, Saliency delivers brief pulses (centre $b \rightarrow 0$, nudge KK , transiently raise TT) to route into **Control** (higher, stable RR , compress) or **Creative** (metastable RR , decompress) and keep the system inside Ψ .

Why this is on-model

- **Griffiths-like bands** emerge naturally on hierarchical/modular networks \rightarrow exactly your Ψ -band premise.
- **Kuramoto on connectomes** shows **power-law event statistics** and **metastability** with inhibitory links/homeostasis—precisely the telemetry you want for band detection and prong-specific states.

In short: **Kuramoto** gives you the clean band + synchrony/metastability lens across the whole network; **Wilson–Cowan** gives you the **E:I knob** and biophysical face of the same controllers. Use them together to instantiate the Trident G-Loop succinctly and testably.

Yes—with this setup you can model **from column** \rightarrow **patch** \rightarrow **system** using the same knobs and telemetry.

How to scale it cleanly

1) Column (micro/meso).

Use **Wilson–Cowan** (E/I) per **microcircuit/column** to sit each unit near a Hopf edge; homeostatic rules keep local **E:I** in the Ψ -window. Read out a local order parameter (e.g., E-activity amplitude/phase). This is where F , $F \cdot F^*$, bb (D1:D2 tilt via pathway gains), and $T(\chi)T(\chi)$ (noise/input variance) act directly.

2) Patch / mesoscale modules.

Group columns into a **hierarchical-modular network** (HMN). Couple their phases (or mean fields) with **Kuramoto-style** interactions; HMNs naturally stretch a single critical point into an

extended critical-like band (Griffiths phase)—i.e., a Ψ -band rather than a knife-edge. Expect power-law episodes and rare-region effects across a **range** of couplings.

3) Whole-brain networks.

Place modules on an empirical **connectome** and evolve a **Kuramoto layer** on top: you get **partial synchrony/metastability** and **power-law desynchronisation durations** with **control-parameter-dependent exponents** below the transition—strong evidence for **extended dynamical criticality** on real graphs. Adding a bit of inhibition or homeostatic gain equalisation preserves these banded signatures.

4) One set of controllers across scales.

- **Thermostat** $F \rightarrow F^*$: slow homeostasis of local gain/coupling toward a target partial synchrony level (keeps units in-band).
- **Precision-bias** b : multiplicatively tilts Control vs Creative subnetworks (e.g., FPN-B \leftrightarrow DMN/FPCN-A) at meso/macro; biases rule-maintenance vs hypothesis search at micro.
- **Temperature** $T(\chi)$: scales noise/frequency jitter when cross-loop inconsistency rises; promotes sampling and re-entry to Ψ .

5) Telemetry invariants (all scales).

- **Band occupancy**: interval with non-saturated synchrony RR and PL event durations (exponent $\tau \sim 1 - \frac{1}{2}$).
- **Metastability peak** near branch point $F \approx F^*$.
- **Rare-region / spectral** markers in HMNs (Lifshitz tails/localised modes). These are exactly what HMN theory and connectome-Kuramoto simulations report.

Why this is empirically reasonable

- HMNs and empirical connectomes **generate** extended critical-like regimes (Griffiths phases) \rightarrow matches your Ψ -band premise.
- Kuramoto on human connectomes shows **smeared crossovers**, **metastability**, and **power-law duration tails** that shift with coupling and inhibition/homeostasis \rightarrow gives you a robust, scalable order parameter.

Quick recipe you can implement

- **Node:** Wilson–Cowan (E/I) \sim column.
- **Edge:** Kuramoto phase coupling between nodes/modules.
- **Controllers:** $F \cdot F^*$ (slow gain homeostasis), bb (subnetwork gain tilt), $T(\chi)T(\chi)$ (noise/jitter).
- **Readouts:** local $E:I$, global/meso $R(t)R(t)$, PL exponents of (de)synchronisation durations, recovery half-life to $F \rightarrow F \cdot F$ to F^* .

That gives you a single, compact formalism that's faithful to **column physiology** and **whole-brain network physics**, while delivering the **Ψ -band + Trident** logic you need.